

### **3. Driver sequence**

NXKR\_YoungSikYang

Exported on 03/12/2024

## Table of Contents

1. Kernel initialization .....	3
1.1 Kernel initialization level.....	3
1.2 Kernel initialization process.....	3
1.2.1 Flow of function calls .....	3
1.2.2 Registration of initcalls .....	4
1.3 Diagram showing how modules are loaded in the kernel initialization.....	7
2. Module driver .....	8
2.1 Registration and loading of a module driver.....	8
2.2 Registering a module driver in 'postcore' .....	9
2.2.1 Check by time .....	9
2.2.2 Check by modifying do_initcalls() .....	9
Reference .....	11

# 1. Kernel initialization

The Linux kernel initialization is a highly structured process that ensures system components and drivers are started in a specific order, allowing dependencies to be resolved correctly.

The initialization sequence consists of several levels.

## 1.1 Kernel initialization level

Each phase is called by `initcall()`

1. **pure**: Architecture-specific, dealing with setting up essential hardware configurations.
2. **core**: Involves setting up essential kernel services and infrastructure, scheduling, interrupt handling, and basic memory management.
3. **postcore**: Initializes additional core subsystems that depend on the very basic services set up during the core phase.
4. **arch**: Since the Linux kernel supports multiple hardware architectures (such as x86, ARM, MIPS, etc.), this phase customizes the initialization process to the specific requirements of the current architecture.
5. **subsys**: The subsystems initialize various kernel subsystems that are not directly tied to the core kernel functionality. This can include driver frameworks, networking, filesystem support, and other subsystems that provide higher-level services.
6. **fs**: The filesystem initialization phase sets up the kernel's filesystem infrastructure, allowing the kernel to access file systems on disk, which is critical for the rest of the system's operation.
7. **rootfs**: The root filesystem is mounted during this phase.
8. **device**: This stage involves the initialization of device drivers and the device model, which allows the kernel to interact with the hardware components of the system.
9. **late**: The late initialization phase is the final step, handling any remaining initialization. This might include starting up user-space applications and services that are essential for the system's operation.

## 1.2 Kernel initialization process

### 1.2.1 Flow of function calls

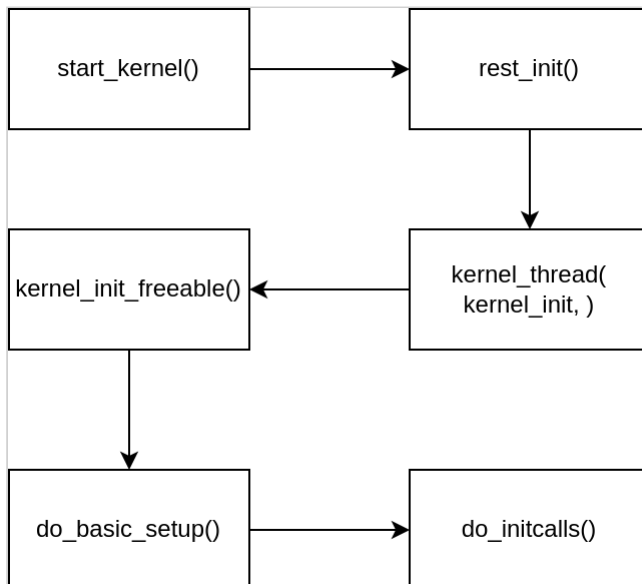
```
head.S runs init/main.c : start_kernel()
```

```

kernel-4.4.x > arch > arm64 > kernel > asm head.S
#include CONFIG_KASLR
    bl kasan_early_init
#endif
#ifdef CONFIG_RANDOMIZE_BASE
    tst x23, ~(MIN_KIMG_ALIGN - 1) // already running randomized?
    b.ne 0f
    mov x0, x21 // pass FDT address in x0
    mov x1, x23 // pass modulo offset in x1
    bl kaslr_early_init // parse FDT for KASLR options
    cbz x0, 0f // KASLR disabled? just proceed
    orr x23, x23, x0 // record KASLR offset
    ret x28 // we must enable KASLR, return
            // to __enable_mmu()
0:
#endif
❖ b start_kernel
ENDPROC(__mmap_switched)

```

**start\_kernel()** initializes various components and leads to **do\_initcalls()**, which runs each init stage.



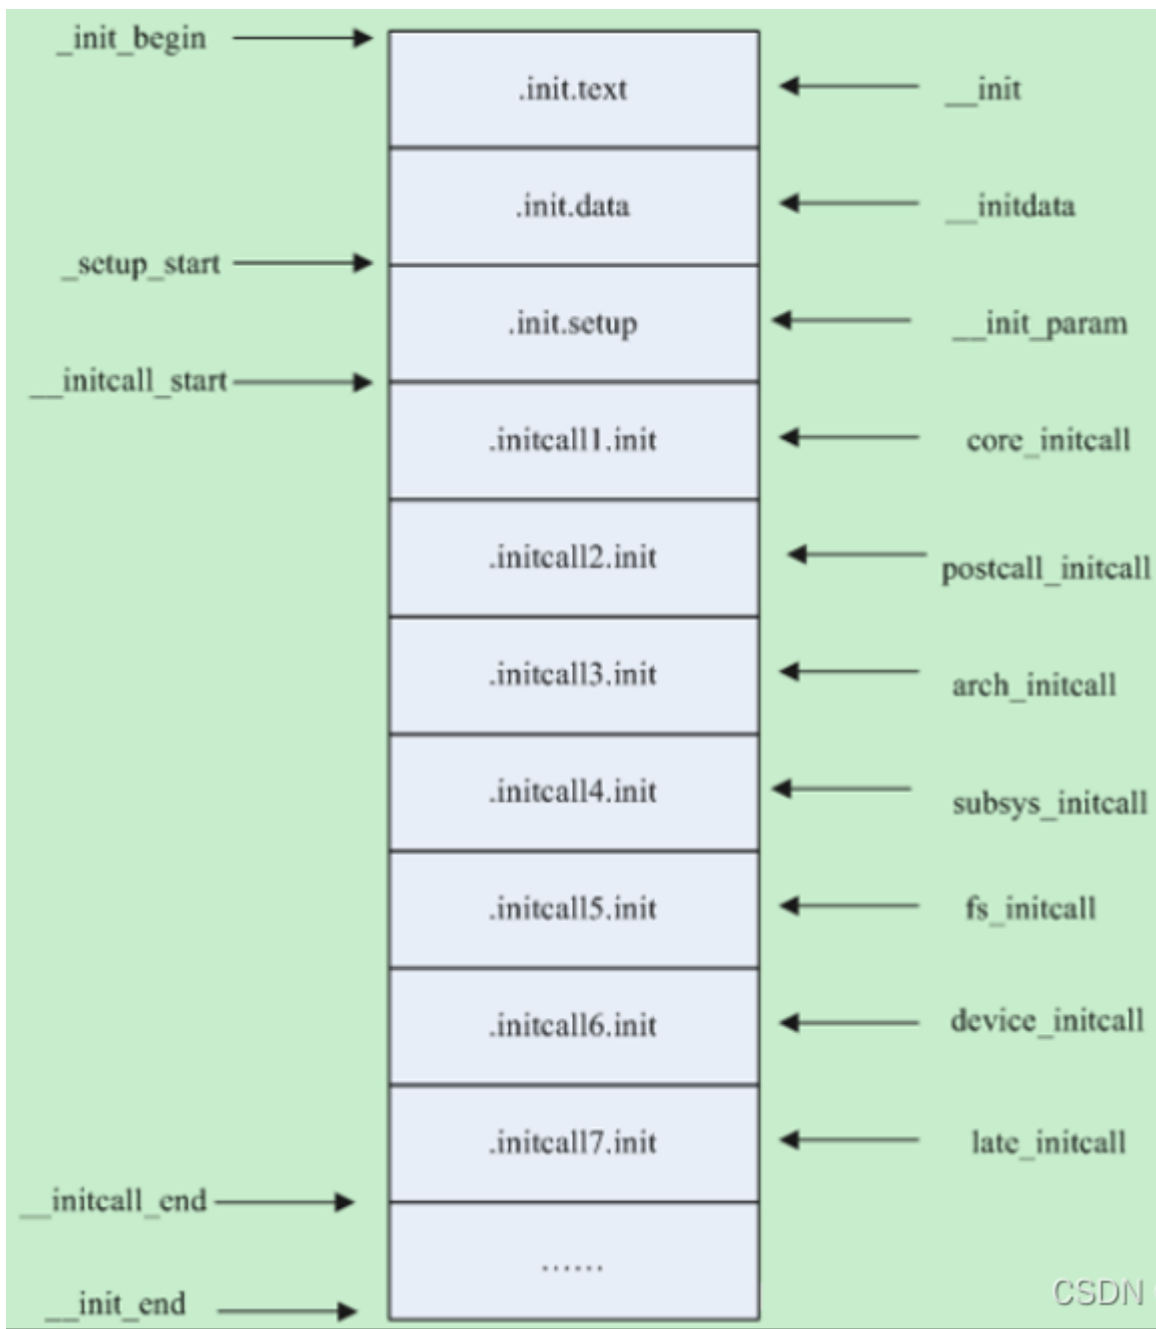
## 1.2.2 Registration of initcalls

Macros defined in `include/linux/init.h` register function pointers in specific sections of the kernel binary so that they can be run in specific init stages.

```
#define pure_initcall(fn)      __define_initcall(fn, 0)

#define core_initcall(fn)     __define_initcall(fn, 1)
#define core_initcall_sync(fn) __define_initcall(fn, 1s)
#define postcore_initcall(fn) __define_initcall(fn, 2)
#define postcore_initcall_sync(fn) __define_initcall(fn, 2s)
#define arch_initcall(fn)     __define_initcall(fn, 3)
#define arch_initcall_sync(fn) __define_initcall(fn, 3s)
#define subsys_initcall(fn)   __define_initcall(fn, 4)
#define subsys_initcall_sync(fn) __define_initcall(fn, 4s)
#define fs_initcall(fn)       __define_initcall(fn, 5)
#define fs_initcall_sync(fn)  __define_initcall(fn, 5s)
#define rootfs_initcall(fn)   __define_initcall(fn, rootfs)
#define device_initcall(fn)   __define_initcall(fn, 6)
#define device_initcall_sync(fn) __define_initcall(fn, 6s)
#define late_initcall(fn)     __define_initcall(fn, 7)
#define late_initcall_sync(fn) __define_initcall(fn, 7s)
```

The linker script(`**vmlinux.lds**`) organizes these special sections so that the function pointers are registered well in their correct section.



Registered functions can be seen using 'nm vmlinux | grep'

```
es/kernel/kernel-4.4.x$ nm vmlinux | grep mydevice_debug_init
ffffff8008add530 d __initcall_mydevice_debug_init2
ffffff8008aad9f4 t mydevice_debug_init
```

During the kernel's boot process, `do_initcalls()` iterates over these sections and executes the function pointers contained within them.

```

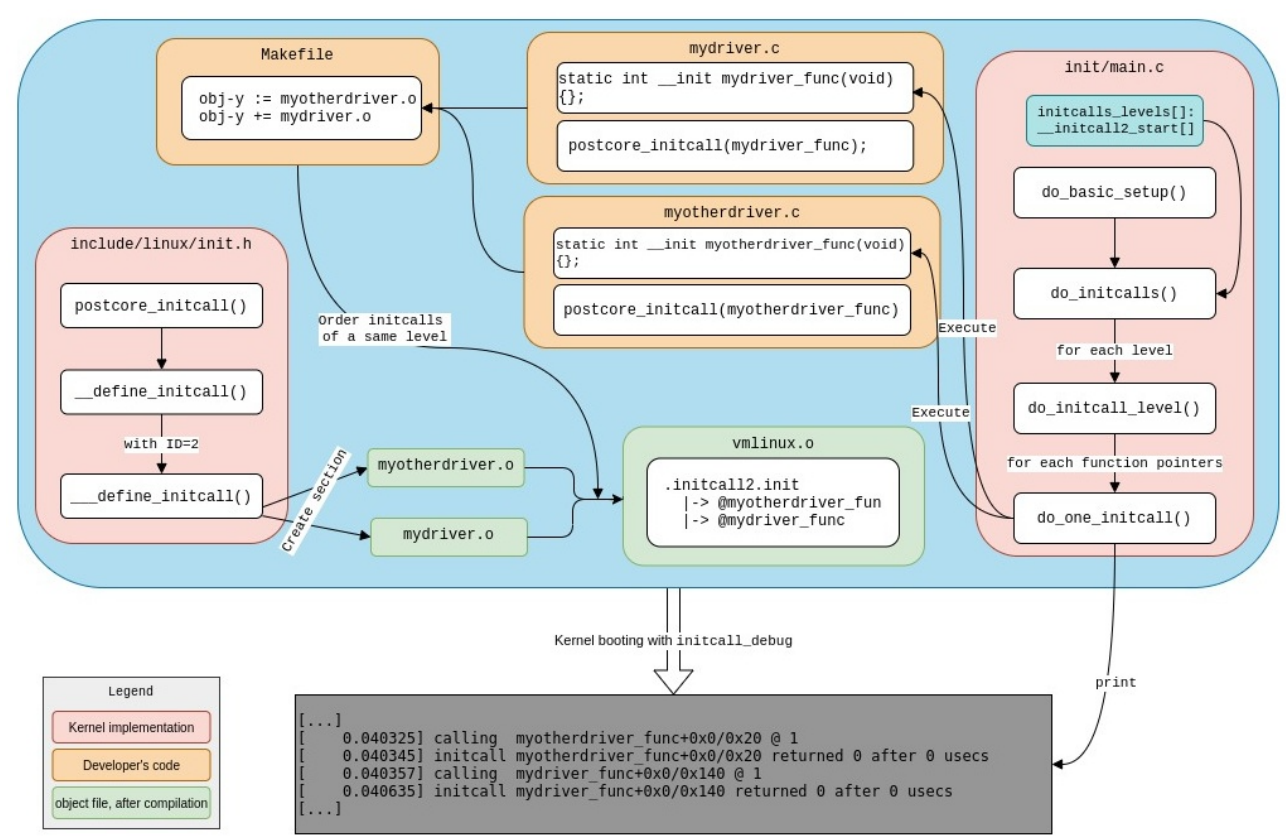
static void init do_initcalls(void)
{
    int level;

#ifdef CONFIG_INITCALLS_THREAD
    for (level = 0; level < CONFIG_INITCALLS_THREAD_LEVEL; level++)
        do_initcall_level(level);

    init_thread_level = level;
    kthread_run(do_initcalls_kth,
        (void *)&init_thread_level, "initcalls:thread");
#else
    for (level = 0; level < ARRAY_SIZE(initcall_levels) - 1; level++)
        do_initcall_level(level);
#endif
}

```

### 1.3 Diagram showing how modules are loaded in the kernel initialization



## 2. Module driver

### 2.1 Registration and loading of a module driver

- `module_init()` registers a module driver

```
static int init misc_init(void)
{
    int error;

    error = misc_register(&tx_misc_device);
    if (error) {
        pr_err("misc_register failed!!!\n");
        return error;
    }

    pr_info("misc_register init done!!!\n");
    return 0;
}
module_init(misc_init)
```

- `module_init()` is defined in `include/linux/module.h`

```
#define module_init(x) __initcall(x);
```

- `__initcall()` is defined in `include/linux/init.h`

```
#define core_initcall(fn)      __define_initcall(fn, 1)
#define core_initcall_sync(fn) __define_initcall(fn, 1s)
#define postcore_initcall(fn) __define_initcall(fn, 2)
#define postcore_initcall_sync(fn) __define_initcall(fn, 2s)
#define arch_initcall(fn)     __define_initcall(fn, 3)
#define arch_initcall_sync(fn) __define_initcall(fn, 3s)
#define subsys_initcall(fn)    __define_initcall(fn, 4)
#define subsys_initcall_sync(fn) __define_initcall(fn, 4s)
#define fs_initcall(fn)        __define_initcall(fn, 5)
#define fs_initcall_sync(fn)    __define_initcall(fn, 5s)
#define rootfs_initcall(fn)     __define_initcall(fn, rootfs)
#define device_initcall(fn)    __define_initcall(fn, 6)
#define device_initcall_sync(fn) __define_initcall(fn, 6s)
#define late_initcall(fn)      __define_initcall(fn, 7)
#define late_initcall_sync(fn) __define_initcall(fn, 7s)

#define __initcall(fn) device_initcall(fn)
```

This shows `module_init()` registers a module driver in the **device** phase.

If it is a built-in module, it is automatically loaded in the **device phase at boot time**.



## 2.2 Registering a module driver in 'postcore'

- Registering the driver in the **late** phase

```
late_initcall(mydevice_debug_init);
```

- Registering the driver in the **postcore** phase

```
postcore_initcall(mydevice_debug_init);
```

### 2.2.1 Check by time

The pictures below show how device driver messages are recorded at different times on different kernel initialization levels.

- Device driver registered in the '**late**' phase using '**late\_initcall()**'.

```
root@s5p6818:~# dmesg | grep yang
[ 2.488000] ===yang's misc device[mydevice_debug_init][L:34]===
[ 2.492000] ===yang's misc device[mydevice_probe][L:19]
```

- Device driver registered in the '**postcore**' phase using '**postcore\_initcall()**'.

```
root@s5p6818:~# dmesg | grep yang
[ 0.184000] ===yang's misc device[mydevice_debug_init][L:34]===
[ 0.208000] ===yang's misc device[mydevice_probe][L:19]
```

The driver init is recorded at an earlier time on **postcore** as shown above.

### 2.2.2 Check by modifying do\_initcalls()

Add a print function to display the current initcall level

```

static void __init do_initcalls(void)
{
    int level;

#ifdef CONFIG_INITCALLS_THREAD
    for (level = 0; level < CONFIG_INITCALLS_THREAD_LEVEL; level++)
        do_initcall_level(level);

    init_thread_level = level;
    kthread_run(do_initcalls_kth,
        (void *)&init_thread_level, "initcalls:thread");
#else
    for (level = 0; level < ARRAY_SIZE(initcall_levels) - 1; level++) {
        pr_info("# Initcall level: %s-----\n", initcall_level_names[level]);
        do_initcall_level(level);
    }
#endif
}

```

You can see the driver is loaded in the **postcore** phase

```

[    0.184000] # Initcall level: postcore-----
[    0.184000] ==yang's misc device[mydevice_debug_init][L:34]==

```

## Reference

- <https://kkslinuxinfo.wordpress.com/2015/12/11/kernel-initialization/>
- [https://blog.csdn.net/weixin\\_47491758/article/details/131157608](https://blog.csdn.net/weixin_47491758/article/details/131157608)
- <https://www.collabora.com/news-and-blog/blog/2020/09/25/initcalls-part-2-digging-into-implementation/>